

Advertising Data Project

In this project we will be working with a fake advertising data set, indicating whether or not a particular internet user clicked on an Advertisement. We will try to create a model that will predict whether or not they will click on an ad based off the features of that user.

This data set contains the following features:

- 'Daily Time Spent on Site': consumer time on site in minutes
- 'Age': customer age in years
- 'Area Income': Avg. Income of geographical area of consumer
- 'Daily Internet Usage': Avg. minutes a day consumer is on the internet
- 'Ad Topic Line': Headline of the advertisement
- 'City': City of consumer
- 'Male': Whether or not consumer was male
- 'Country': Country of consumer
- 'Timestamp': Time at which consumer clicked on Ad or closed window
- 'Clicked on Ad': 0 or 1 indicated clicking on Ad

Let's Get Started

Import necessary Libraries

In [2]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Read "advertising.csv" and set it to dataframe variable

In [3]:

```
df=pd.read_csv("advertising.csv")
```

View the top 5 rows

In [4]:

```
df.head()
```

Out[4]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-03-27 00:53:11
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-04-04 01:39:02
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-03-13 20:35:42
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-01-10 02:31:19
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-06-03 03:36:18

View info of the data

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Daily Time Spent on Site              1000 non-null   float64
1   Age                                    1000 non-null   int64
2   Area Income                           1000 non-null   float64
3   Daily Internet Usage                  1000 non-null   float64
4   Ad Topic Line                         1000 non-null   object
5   City                                   1000 non-null   object
6   Male                                   1000 non-null   int64
7   Country                               1000 non-null   object
8   Timestamp                             1000 non-null   object
9   Clicked on Ad                         1000 non-null   int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

View the basic statistical information about the data

In [6]:

```
df.describe()
```

Out[6]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	65.000200	36.009000	55000.000080	180.000100	0.481000	0.500000
std	15.853615	8.785562	13414.634022	43.902339	0.499889	0.500250
min	32.600000	19.000000	13996.500000	104.780000	0.000000	0.000000
25%	51.360000	29.000000	47031.802500	138.830000	0.000000	0.000000
50%	68.215000	35.000000	57012.300000	183.130000	0.000000	0.500000
75%	78.547500	42.000000	65470.635000	218.792500	1.000000	1.000000
max	91.430000	61.000000	79484.800000	269.960000	1.000000	1.000000

Check for null values

In [7]:

```
df.isna().sum()
```

Out[7]:

Daily Time Spent on Site	0
Age	0
Area Income	0
Daily Internet Usage	0
Ad Topic Line	0
City	0
Male	0
Country	0
Timestamp	0
Clicked on Ad	0
dtype: int64	

View all the countries in our data

In [8]:

```
df.columns
df[['Country']]
```

Out[8]:

	Country
0	Tunisia
1	Nauru
2	San Marino
3	Italy
4	Iceland
...	...
995	Lebanon
996	Bosnia and Herzegovina
997	Mongolia
998	Guatemala
999	Brazil

1000 rows × 1 columns

View all the unique values in 'Ad Topic Line'

In [9]:

```
df["Ad Topic Line"].unique()
```

Out[9]:

```
array(['Cloned 5thgeneration orchestration',
      'Monitored national standardization',
      'Organic bottom-line service-desk',
      'Triple-buffered reciprocal time-frame',
      'Robust logistical utilization', 'Sharable client-driven softw
are',
      'Enhanced dedicated support', 'Reactive local challenge',
      'Configurable coherent function',
      'Mandatory homogeneous architecture',
      'Centralized neutral neural-net',
      'Team-oriented grid-enabled Local Area Network',
      'Centralized content-based focus group',
      'Synergistic fresh-thinking array',
      'Grass-roots coherent extranet',
      'Persistent demand-driven interface',
      'Customizable multi-tasking website', 'Intuitive dynamic attit
ude',
      'Grass-roots solution-oriented conglomeration'.
])
```

View all the cities

In [10]:

```
df[["City"]]
```

Out[10]:

	City
0	Wrightburgh
1	West Jodi
2	Davidton
3	West Terrifurt
4	South Manuel
...	...
995	Duffystad
996	New Darlene
997	South Jessica
998	West Steven
999	Ronniemouth

1000 rows × 1 columns

Change datatype of 'Timestamp' column to datetime format

In []:

In [11]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   Daily Time Spent on Site             1000 non-null   float64
 1   Age                                   1000 non-null   int64   
 2   Area Income                           1000 non-null   float64
 3   Daily Internet Usage                  1000 non-null   float64
 4   Ad Topic Line                         1000 non-null   object  
 5   City                                  1000 non-null   object  
 6   Male                                  1000 non-null   int64   
 7   Country                              1000 non-null   object  
 8   Timestamp                            1000 non-null   object  
 9   Clicked on Ad                         1000 non-null   int64   
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

Create a new column called months by fetching month data from Timestamp

In [12]:

```
df["month"]=pd.DatetimeIndex(df["Timestamp"]).month
df
```

Out[12]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-C 00:3
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-C 01:3
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-C 20:3
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-C 02:3
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-C 03:3
...
995	72.97	30	71384.57	208.58	Fundamental modular algorithm	Duffystad	1	Lebanon	2016-C 21:4
996	51.30	45	67782.17	134.42	Grass-roots cohesive monitoring	New Darlene	1	Bosnia and Herzegovina	2016-C 02:3
997	51.63	51	42415.72	120.37	Expanded intangible solution	South Jessica	1	Mongolia	2016-C 17:2
998	55.55	19	41920.79	187.95	Proactive bandwidth- monitored policy	West Steven	0	Guatemala	2016-C 02:3
999	45.01	26	29875.80	178.35	Virtual 5thgeneration emulation	Ronniemouth	0	Brazil	2016-C 21:4

1000 rows × 11 columns



Create a new column called Year by fetching year from Timestamp

In [13]:

```
df["year"]=pd.DatetimeIndex(df["Timestamp"]).year

df
```

Out[13]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-C 00:00:00
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-C 01:00:00
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-C 20:00:00
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-C 02:00:00
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-C 03:00:00
...
995	72.97	30	71384.57	208.58	Fundamental modular algorithm	Duffystad	1	Lebanon	2016-C 21:00:00
996	51.30	45	67782.17	134.42	Grass-roots cohesive monitoring	New Darlene	1	Bosnia and Herzegovina	2016-C 02:00:00
997	51.63	51	42415.72	120.37	Expanded intangible solution	South Jessica	1	Mongolia	2016-C 17:00:00
998	55.55	19	41920.79	187.95	Proactive bandwidth- monitored policy	West Steven	0	Guatemala	2016-C 02:00:00
999	45.01	26	29875.80	178.35	Virtual 5thgeneration emulation	Ronniemouth	0	Brazil	2016-C 21:00:00

1000 rows × 12 columns



Create an new column called Date by fetching day from Timestamp

In [15]:

```
df["date"]=pd.DatetimeIndex(df["Timestamp"]).day
df
```

Out[15]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-C 00:4
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-C 01:3
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-C 20:3
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-C 02:3
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-C 03:3
...
995	72.97	30	71384.57	208.58	Fundamental modular algorithm	Duffystad	1	Lebanon	2016-C 21:4
996	51.30	45	67782.17	134.42	Grass-roots cohesive monitoring	New Darlene	1	Bosnia and Herzegovina	2016-C 02:0
997	51.63	51	42415.72	120.37	Expanded intangible solution	South Jessica	1	Mongolia	2016-C 17:2
998	55.55	19	41920.79	187.95	Proactive bandwidth- monitored policy	West Steven	0	Guatemala	2016-C 02:3
999	45.01	26	29875.80	178.35	Virtual 5thgeneration emulation	Ronniemouth	0	Brazil	2016-C 21:4

1000 rows × 13 columns

Create a new column called Hour by fetching hour from Timestamp

In [16]:

```
df["hour"]=pd.DatetimeIndex(df["Timestamp"]).hour
df
```

Out[16]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Ad Topic Line	City	Male	Country	Timestamp
0	68.95	35	61833.90	256.09	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	2016-C 00:4
1	80.23	31	68441.85	193.77	Monitored national standardization	West Jodi	1	Nauru	2016-C 01:3
2	69.47	26	59785.94	236.50	Organic bottom-line service-desk	Davidton	0	San Marino	2016-C 20:3
3	74.15	29	54806.18	245.89	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	2016-C 02:3
4	68.37	35	73889.99	225.58	Robust logistical utilization	South Manuel	0	Iceland	2016-C 03:3
...	
995	72.97	30	71384.57	208.58	Fundamental modular algorithm	Duffystad	1	Lebanon	2016-C 21:4
996	51.30	45	67782.17	134.42	Grass-roots cohesive monitoring	New Darlene	1	Bosnia and Herzegovina	2016-C 02:0
997	51.63	51	42415.72	120.37	Expanded intangible solution	South Jessica	1	Mongolia	2016-C 17:2
998	55.55	19	41920.79	187.95	Proactive bandwidth- monitored policy	West Steven	0	Guatemala	2016-C 02:3
999	45.01	26	29875.80	178.35	Virtual 5thgeneration emulation	Ronniemouth	0	Brazil	2016-C 21:4

1000 rows × 14 columns



Visualization


Create a histplot of age

In [17]:

```
sns.histplot(x=df["Age"])
```

Out[17]:

<AxesSubplot: xlabel='Age', ylabel='Count'>




Create a jointplot of Area Income vs Age

In [18]:

```
sns.jointplot(x=df["Age"],y=df["Area Income"])
```

Out[18]:

<seaborn.axisgrid.JointGrid at 0x1e9736e3100>




Create a jointplot showing the kde distributions of Daily Time spend on site vs Age

In [19]:

```
sns.jointplot(x=df["Daily Time Spent on Site"],y=df['Age'],kind='kde')
```

Out[19]:

<seaborn.axisgrid.JointGrid at 0x1e973ad5430>




Create a countplot to show Hour (See which hour the users are most active)

In [144]:

```
sns.countplot(x=df["hour"])
```

Out[144]:

<AxesSubplot:xlabel='hour', ylabel='count'>




Create a countplot to show which day user are the most active

In [22]:

```
sns.countplot(x=df["date"])
```

Out[22]:

<AxesSubplot:xlabel='date', ylabel='count'>



Create a heatmap to visualize the correlation between columns

In [98]:

```
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(),annot=True,fmt=".3f")
```

Out[98]:

<AxesSubplot:>



Split the data into features and target variables (X and y)

In [112]:

```
df = pd.get_dummies(df, drop_first= True)
# Choose the columns you see fit
X=df.drop(columns="Clicked on Ad")
```

In [113]:

```
y=df["Clicked on Ad"]
```

Standardize the data


In [114]:

```
from sklearn.preprocessing import StandardScaler
df = pd.get_dummies(df, drop_first=True)
df.head()
```

Out[114]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Male	Clicked on Ad	month	year	hour	Ad Topic Line_Adaptive asynchronous attitude	...	Ti
0	68.95	35	61833.90	256.09	0	0	3	2016	0	0	...	
1	80.23	31	68441.85	193.77	1	0	4	2016	1	0	...	
2	69.47	26	59785.94	236.50	0	0	3	2016	20	0	...	
3	74.15	29	54806.18	245.89	1	0	1	2016	2	0	...	
4	68.37	35	73889.99	225.58	0	0	6	2016	3	0	...	

5 rows × 3211 columns



In [115]:

```
st=StandardScaler()
xcolumns=X.columns
st.fit_transform(X)

k=pd.DataFrame(df,columns=xcolumns)
```


In [116]:

```
df.head()
```

Out[116]:

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	Male	Clicked on Ad	month	year	hour	Ad Topic Line_Adaptive asynchronous attitude	...	Ti
0	68.95	35	61833.90	256.09	0	0	3	2016	0	0	...	
1	80.23	31	68441.85	193.77	1	0	4	2016	1	0	...	
2	69.47	26	59785.94	236.50	0	0	3	2016	20	0	...	
3	74.15	29	54806.18	245.89	1	0	1	2016	2	0	...	
4	68.37	35	73889.99	225.58	0	0	6	2016	3	0	...	

5 rows × 3211 columns

Split the data into training and testing set

In [118]:

```
from sklearn.model_selection import train_test_split
```

In [119]:

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=100)
```

Create a Logistic Regression model and train it

In [121]:

```
from sklearn.linear_model import LogisticRegression
```

In [122]:

```
lr=LogisticRegression()
```

In [123]:

```
#train the model
lr.fit(X_train,y_train)
```

Out[123]:

```
LogisticRegression()
```

Check the accuracy of our model

In [124]:

```
lr.score(X_train,y_train)
```

Out[124]:

```
0.9742857142857143
```

Make prediction using the X_test

In [129]:

```
y_pred=lr.predict(X_test)  
y_pred
```

Out[129]:

```
array([1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1,  
       0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0,  
       0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,  
       0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1,  
       1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1,  
       0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1,  
       1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0,  
       1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1,  
       1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1,  
       1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0,  
       0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0,  
       1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0,  
       1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0,  
       0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1], dtype=int64)
```

Check how accurate the prediction is:

In [139]:

```
from sklearn import metrics
```

In [141]:

```
metrics.mean_squared_error(y_test,y_pred)
```

Out[141]:

```
0.03
```

Check the confusion matrix

In [89]:


```
metrics.confusion_matrix(y_test, y_pred)
```

```
[[145  0]
 [ 9 146]]
```

Plot the confusion matrix on a heatmap

In [137]:

```
sns.heatmap(metrics.confusion_matrix(y_test, y_pred), annot = True)
plt.show()
```



Create a classification report for the model

In [138]:

```
print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	145
1	1.00	0.94	0.97	155
accuracy			0.97	300
macro avg	0.97	0.97	0.97	300
weighted avg	0.97	0.97	0.97	300

Great Job!